

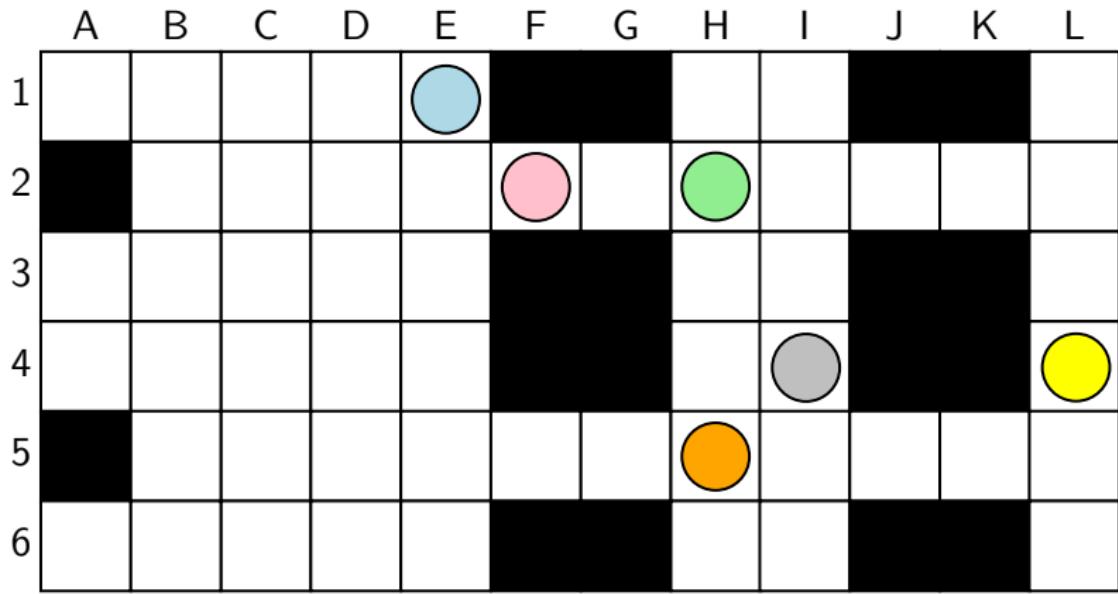
# Multi-Agent Path Finding

Bojie Shen  
Monash University

<https://bojie-shen.com/>

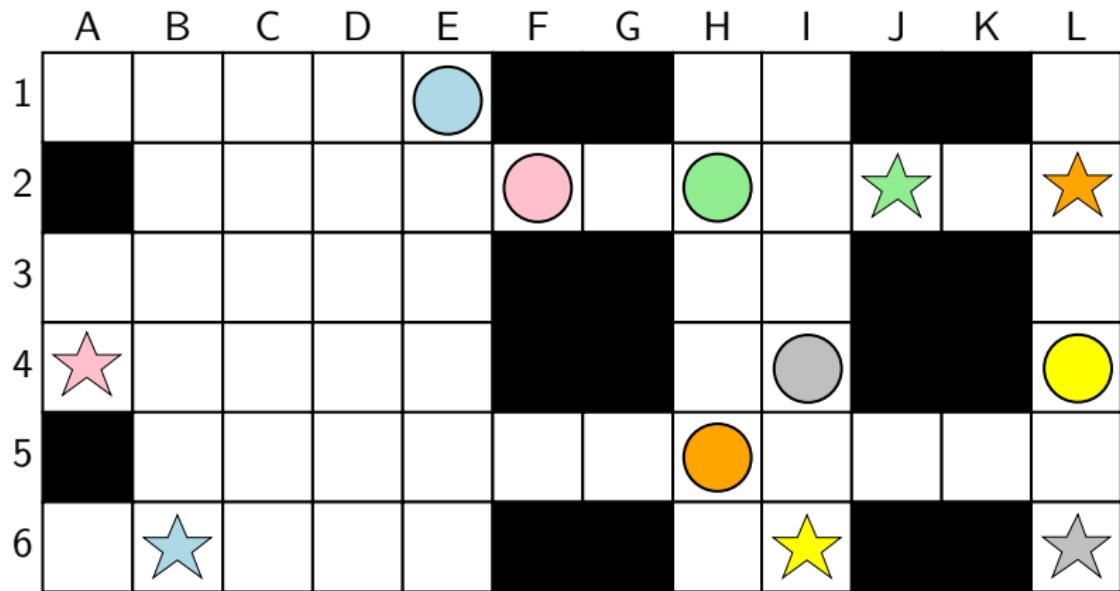
## Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



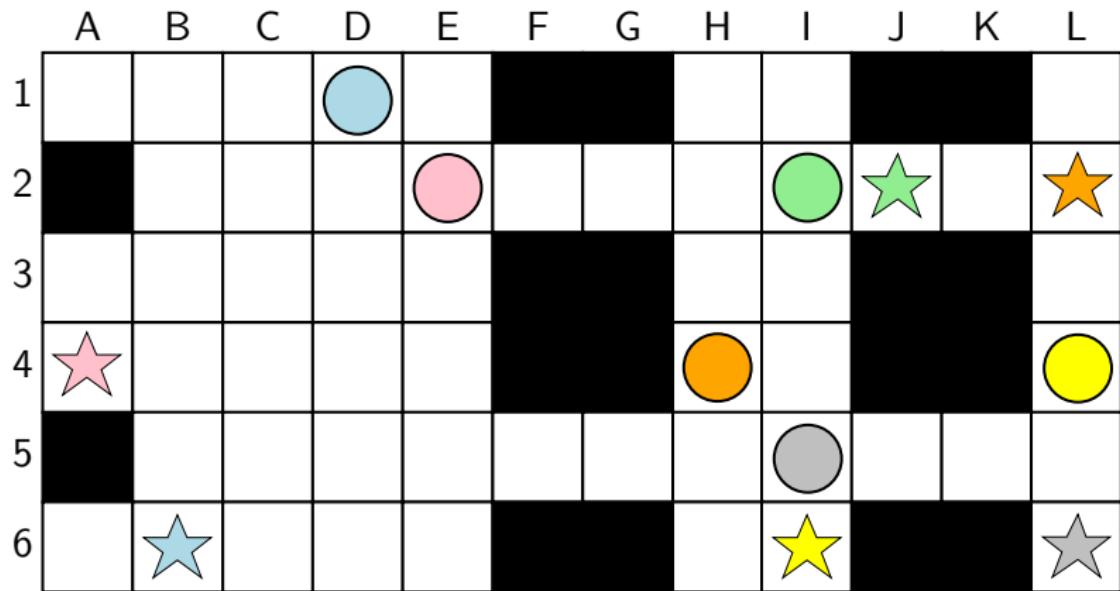
# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



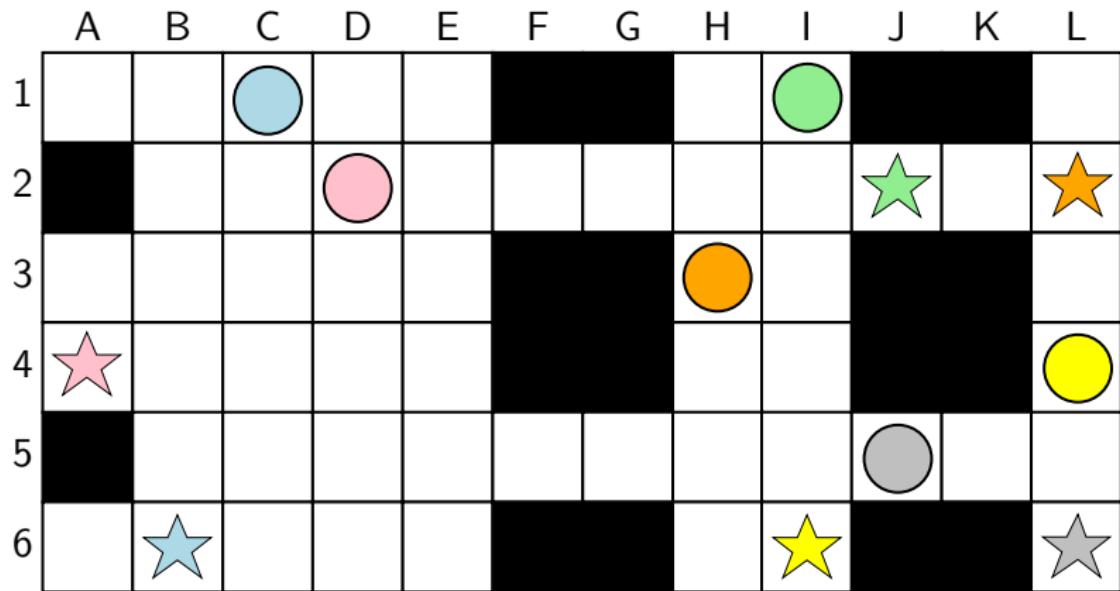
# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



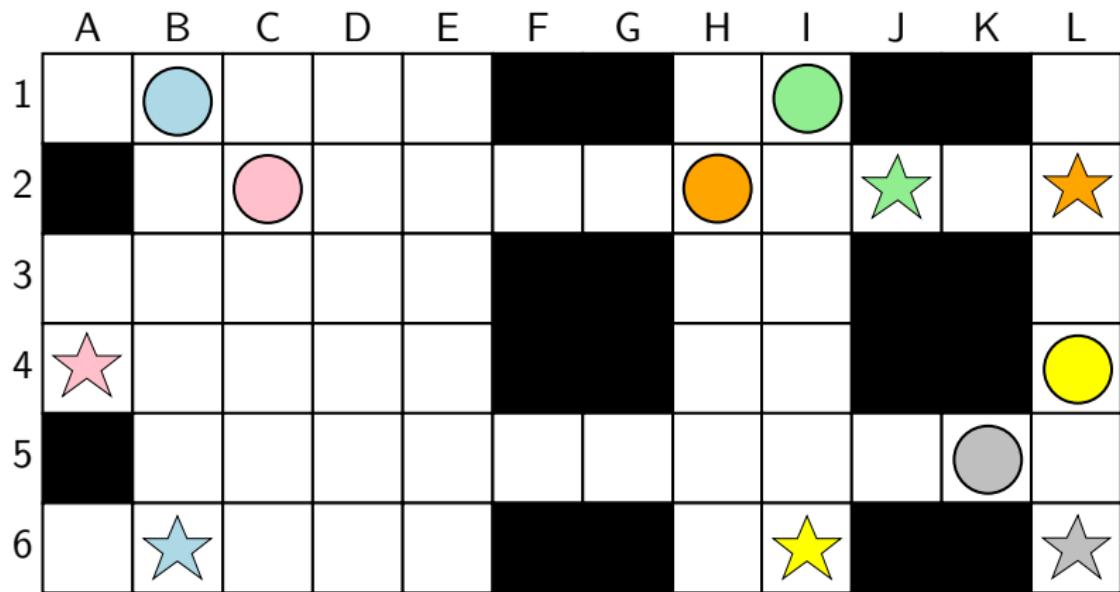
# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



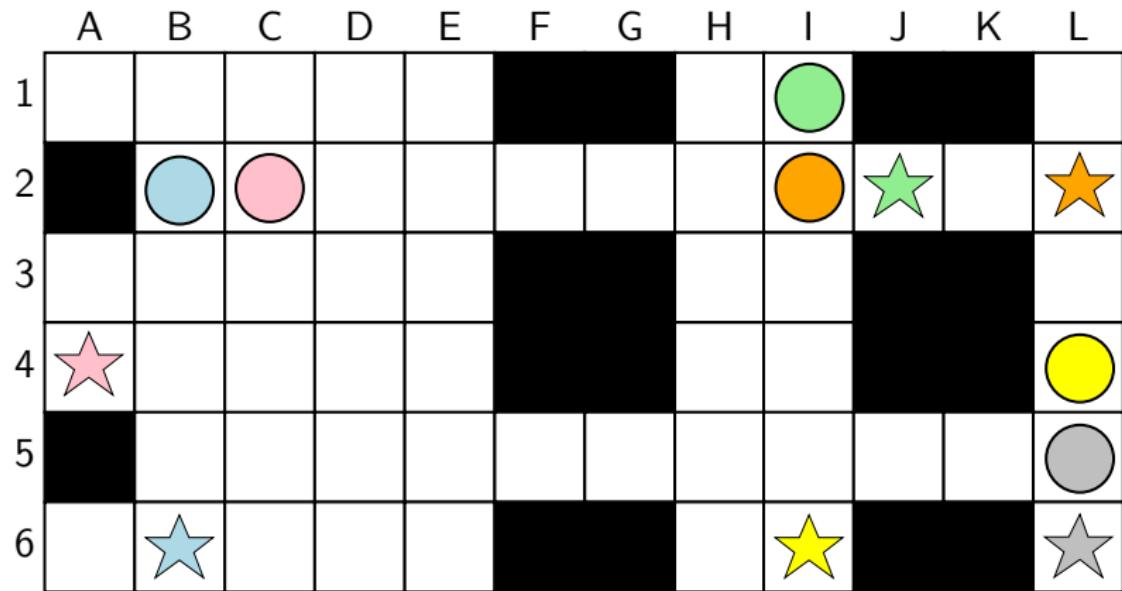
# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



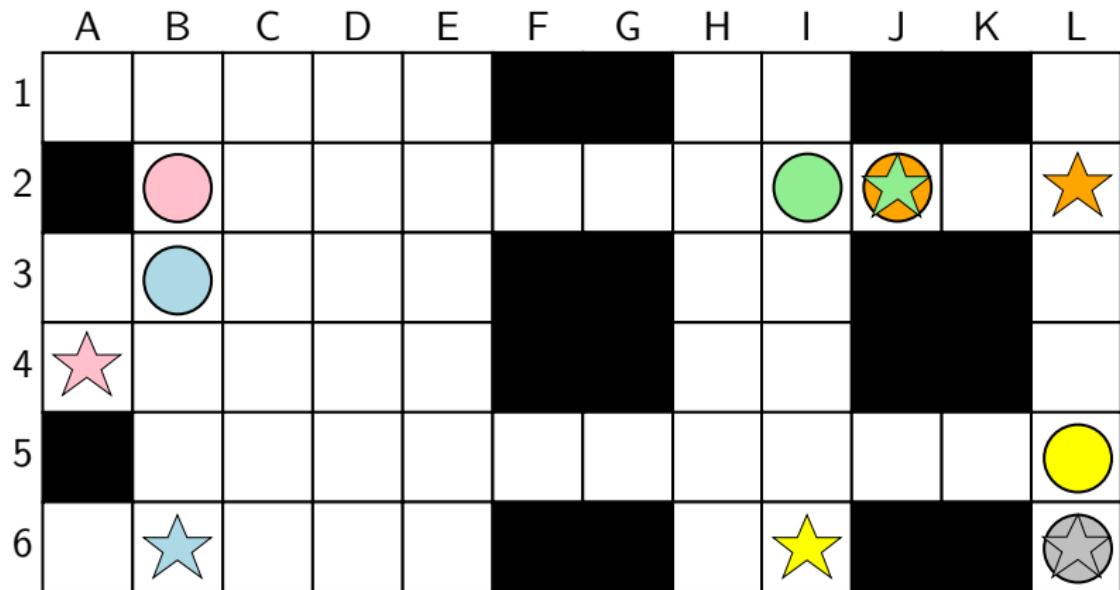
# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



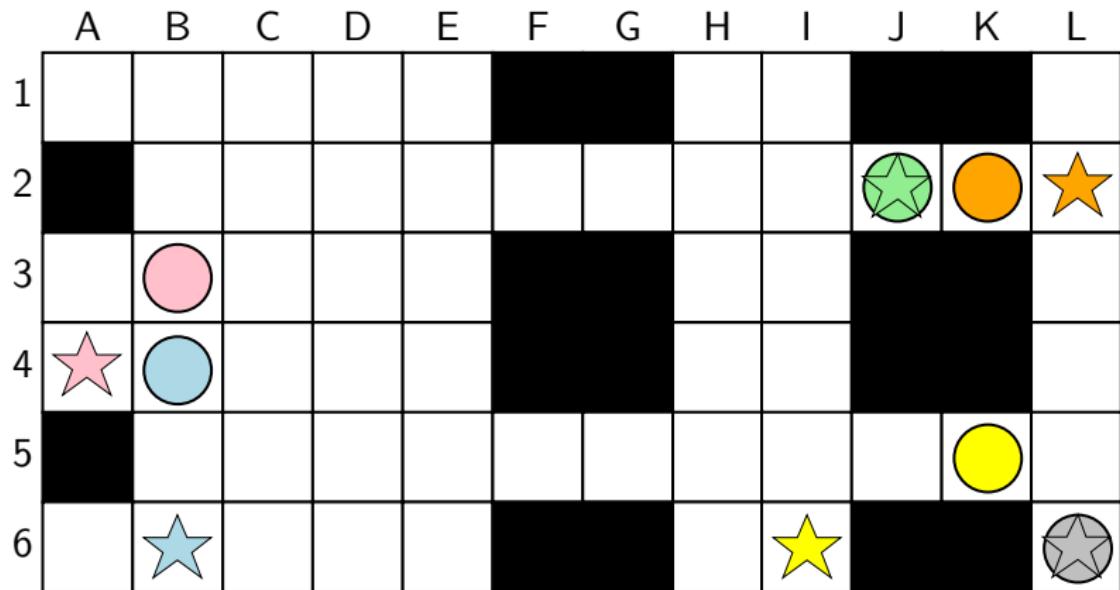
# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



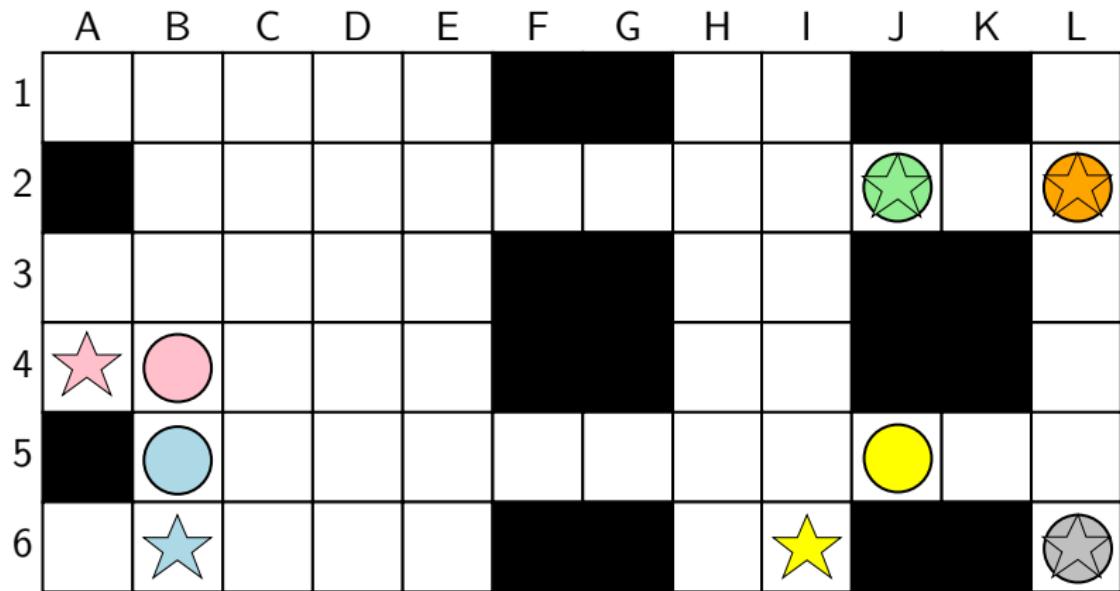
# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



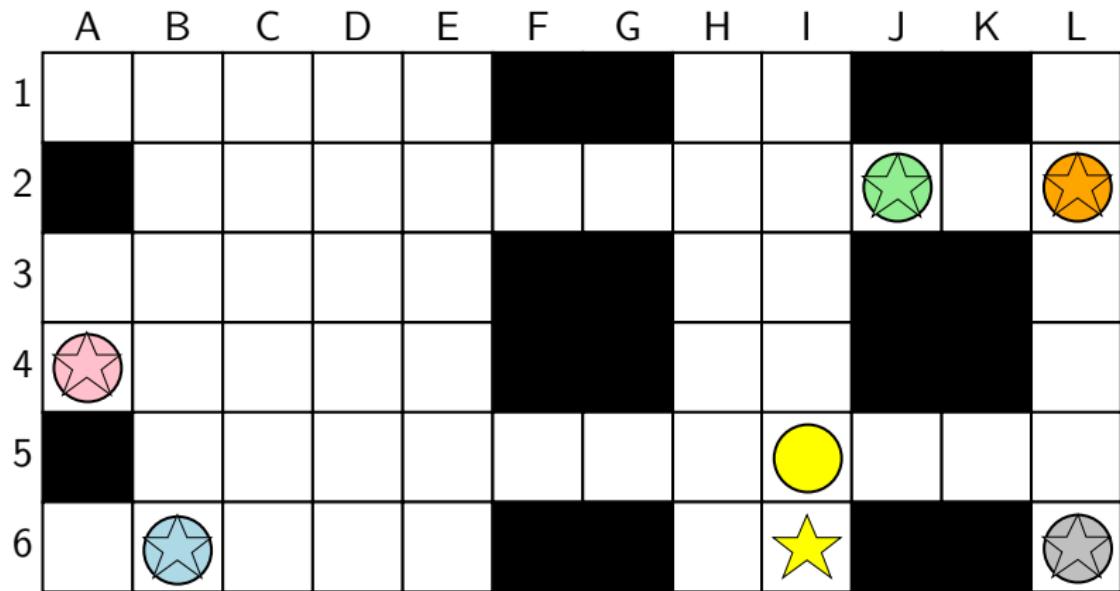
# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



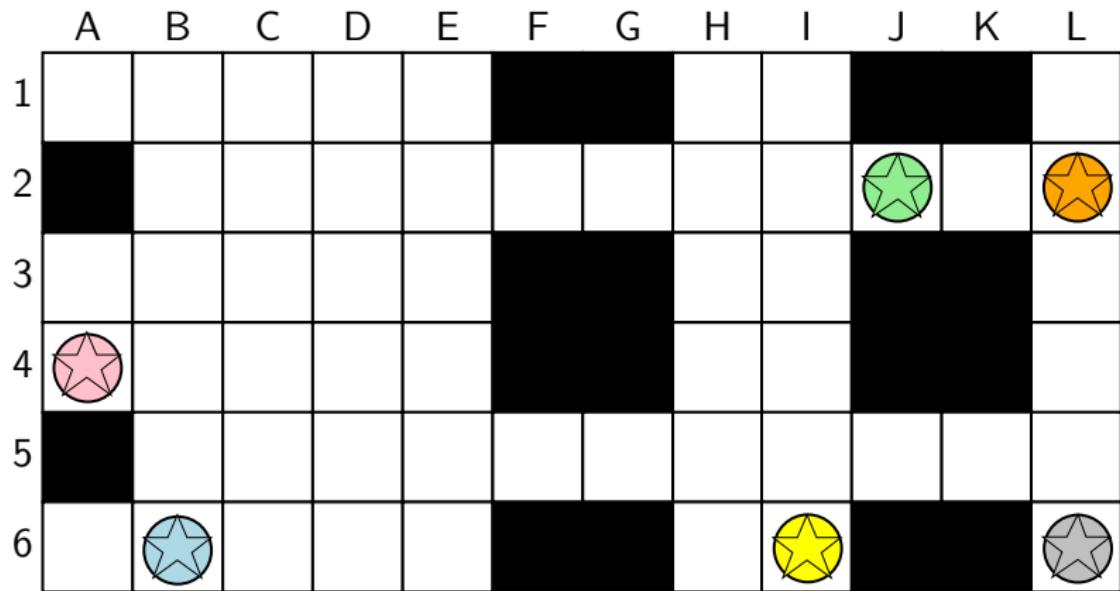
# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



# Recap

The Multi-agent Path Finding problem (MAPF) asks for a collision-free plan that moves a team of agents.



# Classical MAPF

## Setup

- ▶ Each agent begins at a specific start position
- ▶ Each agent finishes at a specific target position
- ▶ All agents move on a 4-connected grid
- ▶ All actions incur the same unit cost
- ▶ Waiting is allowed

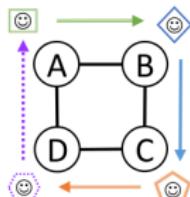
## Objectives

- ▶ **Sum of Costs:** minimise the total cost of all assigned paths.
- ▶ **Makespan:** minimise the arrival time of the last agent.

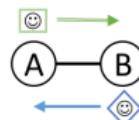
# Classical MAPF

## Constraints

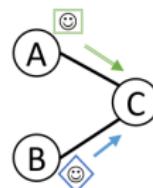
- ▶ Two agents cannot occupy a grid cell at the same time
- ▶ Two agents cannot swap positions at the same time
- ▶ Agents **can** move in “lockstep”
- ▶ Agents **arrive** their destination if they can remain there collision-free



Lockstep movement



Edge conflict



Vertex conflict

# MAPF Applications



Robocup Rescue

## MAPF Applications



## StarCraft II

# MAPF Applications



Amazon Warehouse

# MAPF in Practice

Applications of MAPF often have many additional features. But the core coordination problem is always there.

	<b>MAPF Problem</b>	<b>Amazon Problem</b>
<b>Environment</b>	Grid	Grid
<b>Workload</b>	One shot	Lifelong
<b>Kinematics</b>	No	Yes
<b>Objective</b>	Sum-of-Costs	\_(_ツ)_/`
<b>Solutions</b>	Optimal	Best available

# MAPF approaches

## Centralised

The agents are planned together by a single omniscient controller.

- ▶ Coupled: agents are planned **jointly**.
- ▶ Decoupled: agents are planned **independently**, subject to constraints.

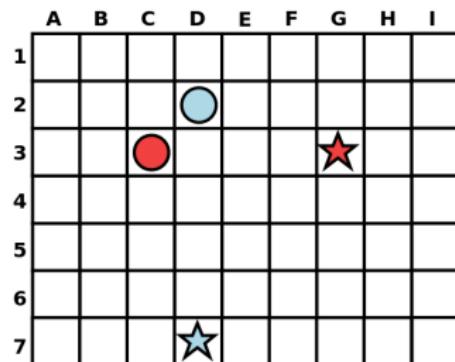
## Distributed

The agents are planned independently, without global knowledge.

- ▶ Each agent has its own model of the task environment
- ▶ Each agent plans for itself and starts executing
- ▶ Conflicts are detected during execution and agents replan
- ▶ (Optional) Local information sharing and local negotiation

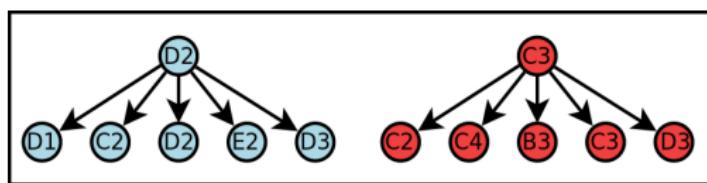
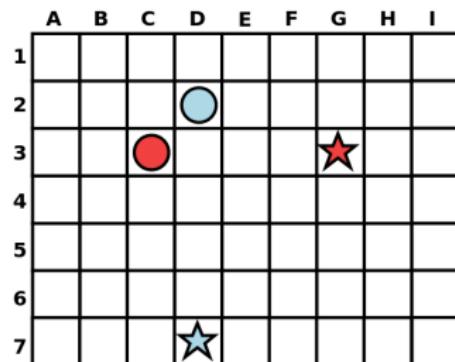
# Centralised A\*

**Idea:** Joint-space planning. We treat the group of agents,  $A$ , as one large agent with many degrees-of-freedom.



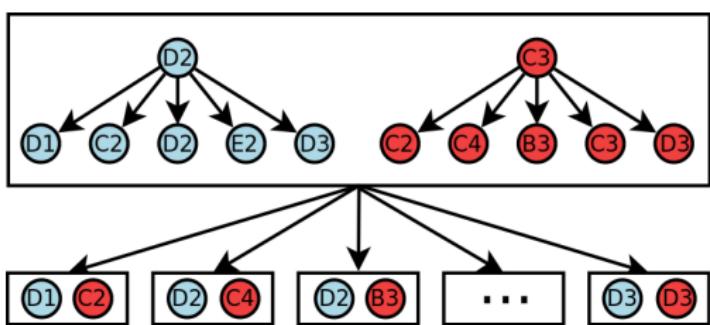
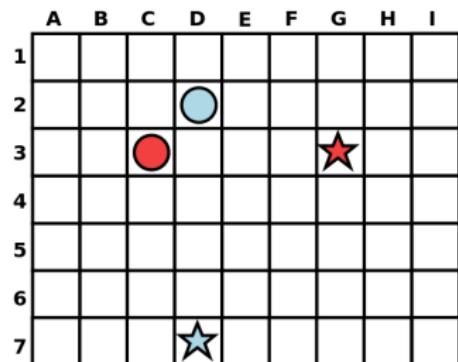
# Centralised A\*

**Idea:** Joint-space planning. We treat the group of agents,  $A$ , as one large agent with many degrees-of-freedom.



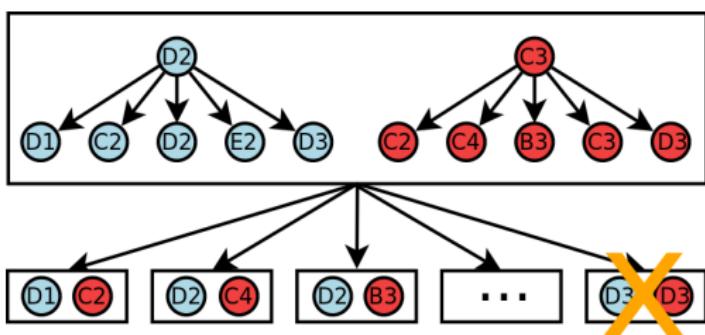
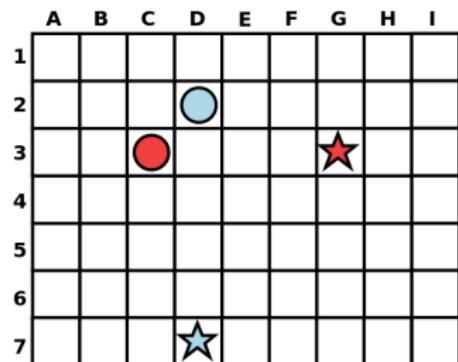
# Centralised A\*

**Idea:** Joint-space planning. We treat the group of agents,  $A$ , as one large agent with many degrees-of-freedom.



# Centralised A\*

**Idea:** Joint-space planning. We treat the group of agents,  $A$ , as one large agent with many degrees-of-freedom.



# Analysis of Centralised A\*

Joint-space planning is complete and optimal. But it doesn't scale and is often a non-starter.

# Analysis of Centralised A\*

Joint-space planning is complete and optimal. But it doesn't scale and is often a non-starter.

Branching factor is  $b^{|A|}$ . On 4-C gridmaps we have  $b = 5$ .

- ▶ 2 agents =  $5^2$  successors (= 25; acceptable)

# Analysis of Centralised A\*

Joint-space planning is complete and optimal. But it doesn't scale and is often a non-starter.

Branching factor is  $b^{|A|}$ . On 4-C gridmaps we have  $b = 5$ .

- ▶ 2 agents =  $5^2$  successors (= 25; acceptable)
- ▶ 3 agents =  $5^3$  successors (=125; challenging)

# Analysis of Centralised A\*

Joint-space planning is complete and optimal. But it doesn't scale and is often a non-starter.

Branching factor is  $b^{|A|}$ . On 4-C gridmaps we have  $b = 5$ .

- ▶ 2 agents =  $5^2$  successors (= 25; acceptable)
- ▶ 3 agents =  $5^3$  successors (=125; challenging)
- ▶ 5 agents =  $5^5$  successors (=3125; unreasonable)

# Analysis of Centralised A\*

Joint-space planning is complete and optimal. But it doesn't scale and is often a non-starter.

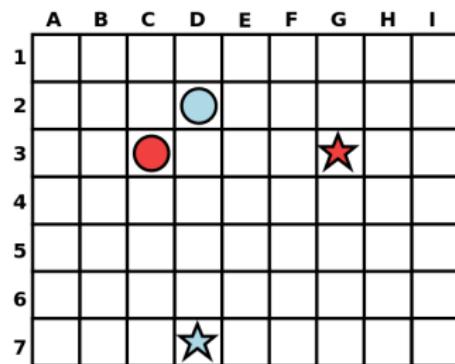
Branching factor is  $b^{|A|}$ . On 4-C gridmaps we have  $b = 5$ .

- ▶ 2 agents =  $5^2$  successors (= 25; acceptable)
- ▶ 3 agents =  $5^3$  successors (=125; challenging)
- ▶ 5 agents =  $5^5$  successors (=3125; unreasonable)
- ▶ 10 agents =  $5^{10}$  (= **lolno**)

CA\* performance can be improved using *partial expansion*. This strategy expands a node multiple times, each time generating only the most promising successors from those that remain. This is a very useful idea and generally applicable. For details see (Felner, A. et al. 2012. Partial-Expansion A\* with Selective Node Generation. In Proceedings of AAAI (pp. 180-181))

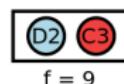
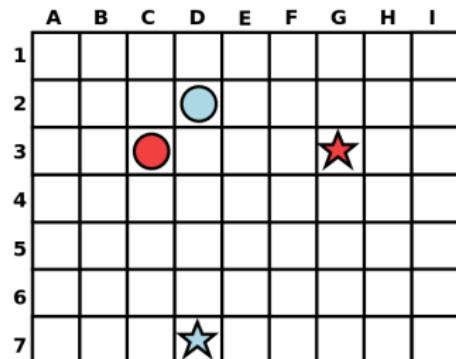
# Operator Decomposition

**Idea:** Agents select compatible moves one at a time, in some fixed order.  
After  $|A|$  individual moves, the search advances one timestep.



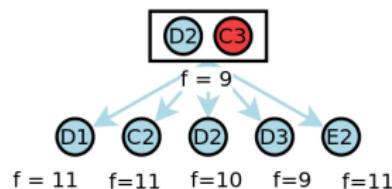
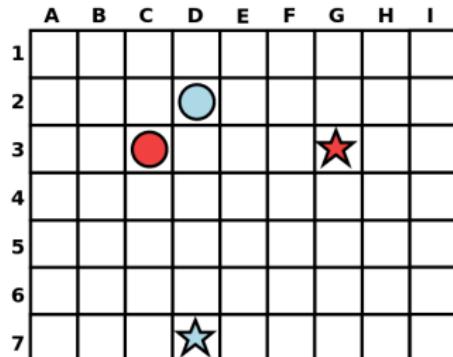
# Operator Decomposition

**Idea:** Agents select compatible moves one at a time, in some fixed order.  
After  $|A|$  individual moves, the search advances one timestep.



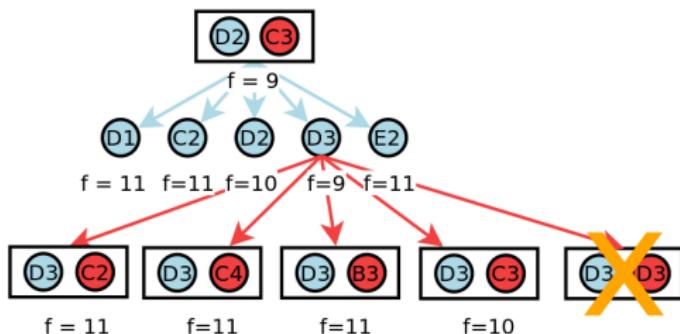
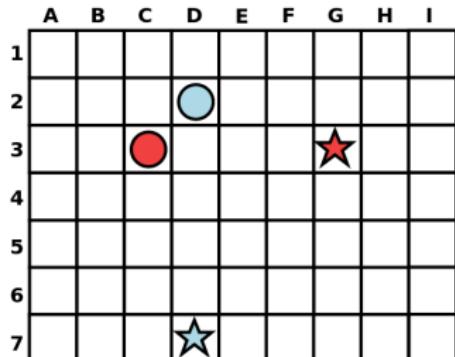
# Operator Decomposition

**Idea:** Agents select compatible moves one at a time, in some fixed order. After  $|A|$  individual moves, the search advances one timestep.



# Operator Decomposition

**Idea:** Agents select compatible moves one at a time, in some fixed order. After  $|A|$  individual moves, the search advances one timestep.



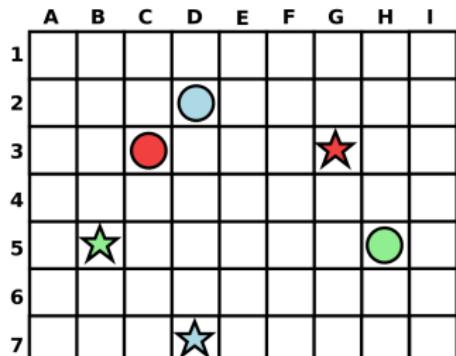
# Analysis of OD

Operator Decomposition is complete and optimal.

- ▶ Branching factor reduced to  $b$  ( $=5$  on 4-C gridmaps)
- ▶ But solutions are deeper in the tree (from  $d$  with CA\* to  $d \times |A|$ )
- ▶ Works well for problems with small numbers of agents

# Independence Detection (OD+ID)

**Idea:** many agents don't interact with other agents or only just a few.  
Try to identify such groups and solve them independently.



---

## Algorithm 1 Simple Independence Detection

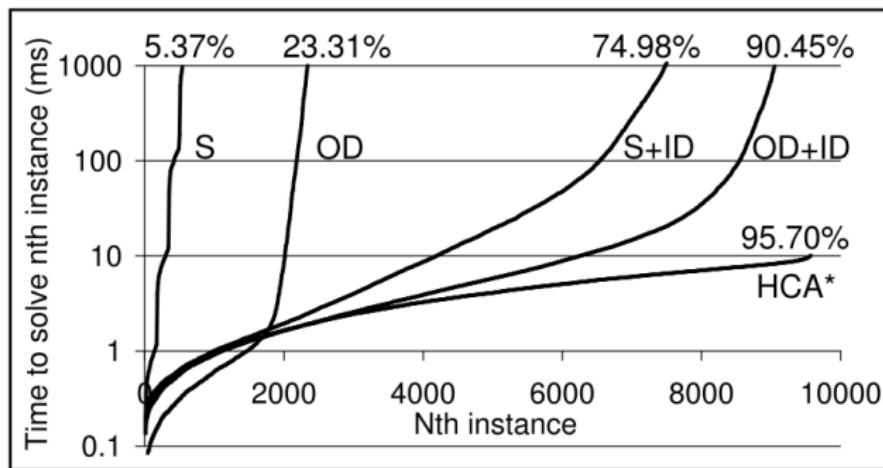
---

- 1: assign each agent to a singleton group
  - 2: plan a path for each group
  - 3: **repeat**
  - 4:   simulate execution of all paths until a conflict occurs
  - 5:   merge two conflicting groups into a single group
  - 6:   cooperatively plan new group
  - 7: **until** no conflicts occur
  - 8: *solution*  $\leftarrow$  paths of all groups combined
  - 9: **return** *solution*
- 

OD+ID is easy to understand and, even 10 years since its introduction, still reasonably effective. It first appears in: (Standley, T.S., 2010, July. Finding Optimal Solutions to Cooperative Pathfinding Problems. In AAAI (Vol. 1, pp. 28-29)).

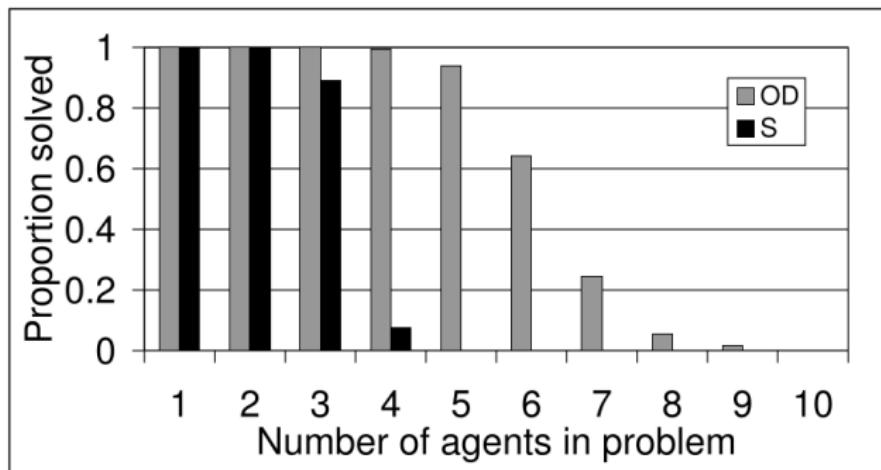
# Effectiveness of OD+ID

Experiments on  $32 \times 32$  grids with 20% random obstacles.  
Random problems with up to 60 agents.

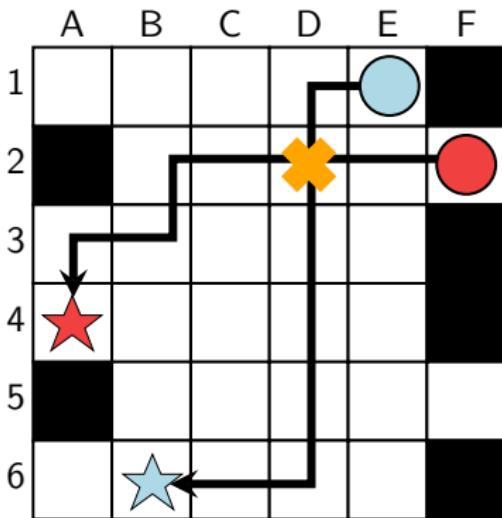


# Effectiveness of OD+ID

Experiments on  $32 \times 32$  grids with 20% random obstacles.  
Random problems with up to 60 agents.

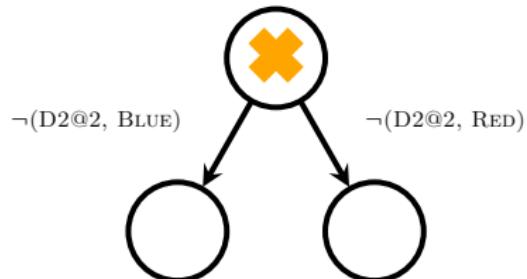
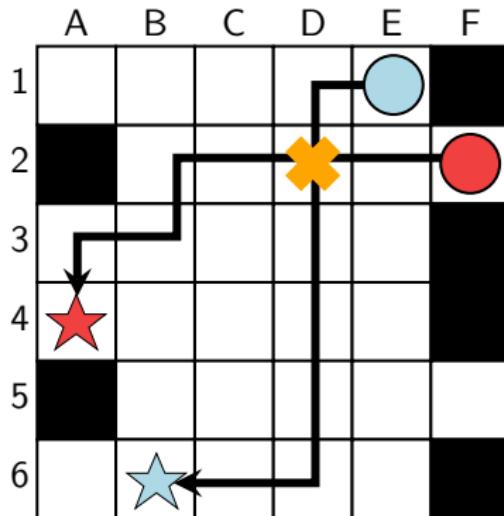


# Conflict-based Search [Sharon et al, AIJ 2015]



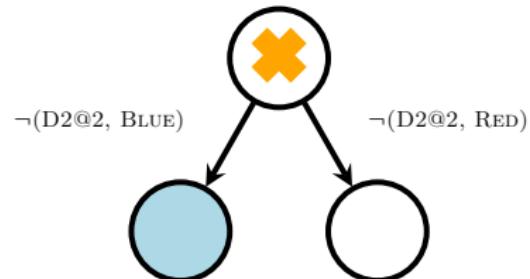
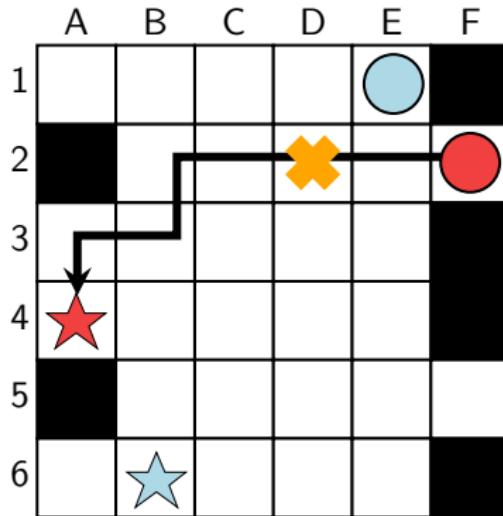
CBS is an optimal branch-and-replan search algorithm.

# Conflict-based Search [Sharon et al, AIJ 2015]



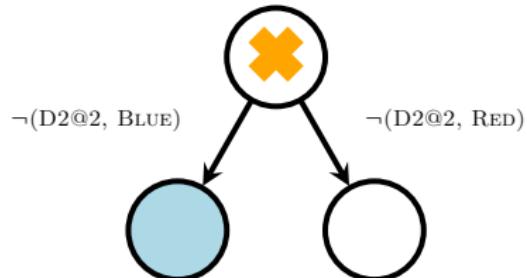
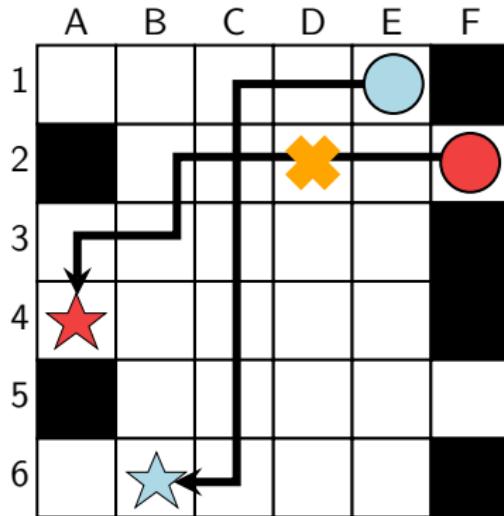
CBS is an optimal branch-and-replan search algorithm.

# Conflict-based Search [Sharon et al, AIJ 2015]



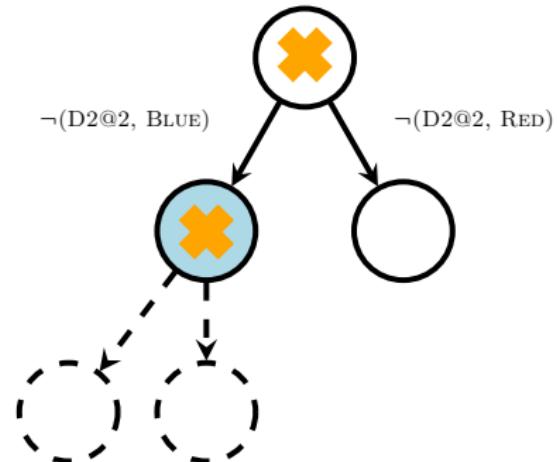
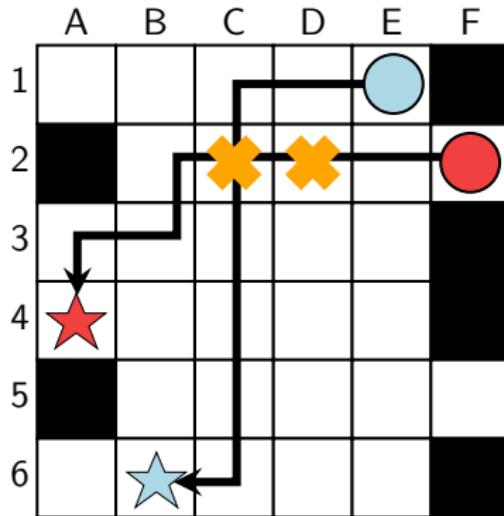
CBS is an optimal branch-and-replan search algorithm.

# Conflict-based Search [Sharon et al, AIJ 2015]



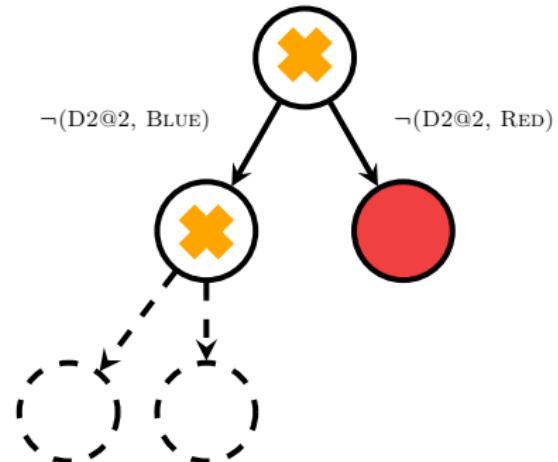
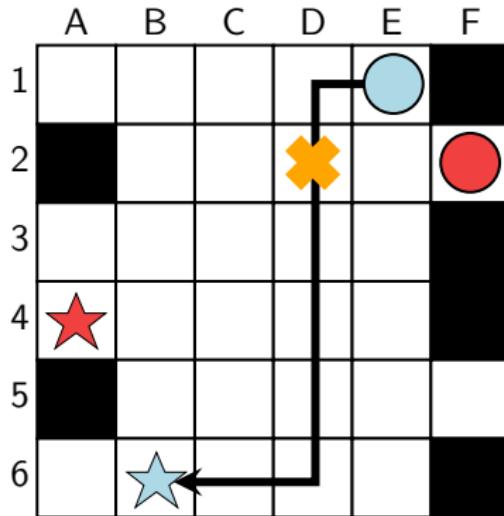
CBS is an optimal branch-and-replan search algorithm.

# Conflict-based Search [Sharon et al, AIJ 2015]



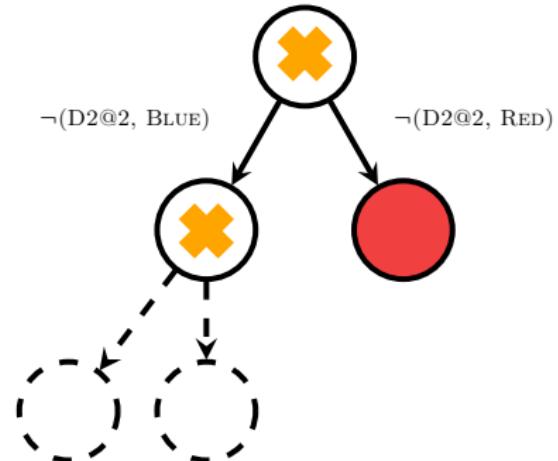
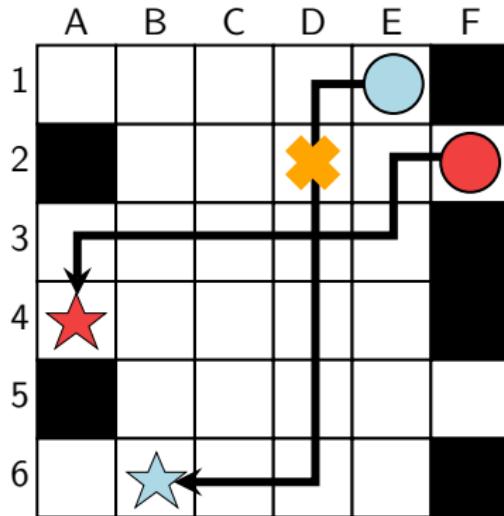
CBS is an optimal branch-and-replan search algorithm.

# Conflict-based Search [Sharon et al, AIJ 2015]



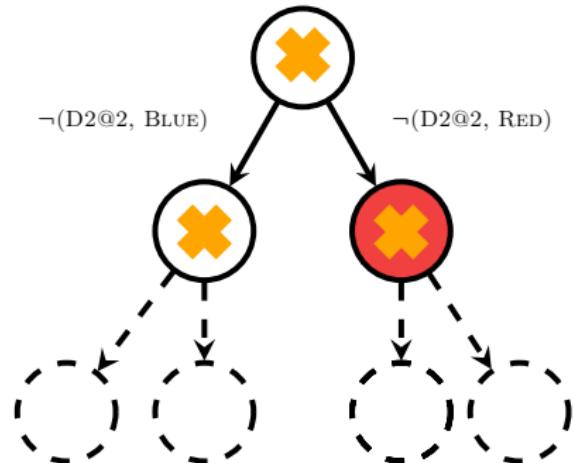
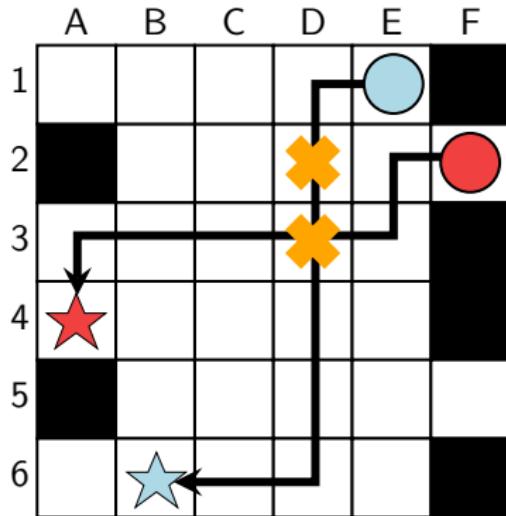
CBS is an optimal branch-and-replan search algorithm.

# Conflict-based Search [Sharon et al, AIJ 2015]



CBS is an optimal branch-and-replan search algorithm.

# Conflict-based Search [Sharon et al, AIJ 2015]



CBS is an optimal branch-and-replan search algorithm.

# Analysis of CBS

CBS is solution complete and optimal.

Optimality Sketch:

- ▶ For each valid solution  $S$ , there exists a node  $N$  on OPEN which **permits**  $S$ .

# Analysis of CBS

CBS is solution complete and optimal.

Optimality Sketch:

- ▶ For each valid solution  $S$ , there exists a node  $N$  on OPEN which **permits**  $S$ .
  - ▶ Trivially true at the root node (no constraints)
  - ▶ Each split divides the set of valid solutions at the parent node.
  - ▶ Constraints eliminate valid solutions:  
from one subtree or the other but never from both.

# Analysis of CBS

CBS is solution complete and optimal.

Optimality Sketch:

- ▶ For each valid solution  $S$ , there exists a node  $N$  on OPEN which **permits**  $S$ .
  - ▶ Trivially true at the root node (no constraints)
  - ▶ Each split divides the set of valid solutions at the parent node.
  - ▶ Constraints eliminate valid solutions:  
from one subtree or the other but never from both.
- ▶ Each node  $N$  is a minimum cost plan where the paths of the agents are consistent with the set of constraints at  $N$

# Analysis of CBS

CBS is solution complete and optimal.

Optimality Sketch:

- ▶ For each valid solution  $S$ , there exists a node  $N$  on OPEN which **permits**  $S$ .
  - ▶ Trivially true at the root node (no constraints)
  - ▶ Each split divides the set of valid solutions at the parent node.
  - ▶ Constraints eliminate valid solutions:  
from one subtree or the other but never from both.
- ▶ Each node  $N$  is a minimum cost plan where the paths of the agents are consistent with the set of constraints at  $N$
- ▶ The minimum f-value of any node  $N \in \text{OPEN}$  is lower-bound on the minimum cost of any valid solution. □

# Analysis of CBS

CBS is solution complete and optimal.

Completeness sketch:

- ▶ After some maximum timestep  $T$  no further collisions are possible (every agent arrives)
- ▶ The cost of every generated node is monotonically non-decreasing.
- ▶ The optimal solution cost  $C \leq T$  □

# Complexity of MAPF

MAPF is intractable in a large number of cases. But not always!

Unlikely to find polynomial time algorithms for these problems:

- ▶ MAPF on general directed graphs
- ▶ MAPF on general undirected graphs
- ▶ MAPF on general grid graphs

MAPF complexity proofs are often very creative and they provide us with wonderful insights into the core difficulties. The following recent paper is a good start for those who want to know more: (**Nebel, B., 2020. On the Computational Complexity of Multi-Agent Pathfinding on Directed Graphs. In Proceedings of ICAPS.**)

# Complexity of MAPF

MAPF is intractable in a large number of cases. But not always!

Cases which admit polynomial-time solutions:

- ▶ Bi-connected graphs
- ▶ Well-formed instances

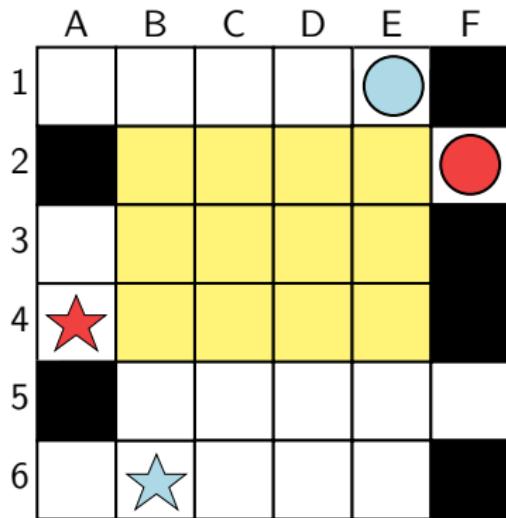
The class of well-formed MAPF problems is proving especially important for practical settings. The following paper is a good starting point for this important topic: **(Cap, M., Vokrinek, J. and Kleiner, A. 2015. Complete decentralized method for on-line multi-robot trajectory planning in well-formed infrastructures. In Proceedings of ICAPS)**

# Multi-Agent Path Finding: Constraints

Bojie Shen  
Monash University

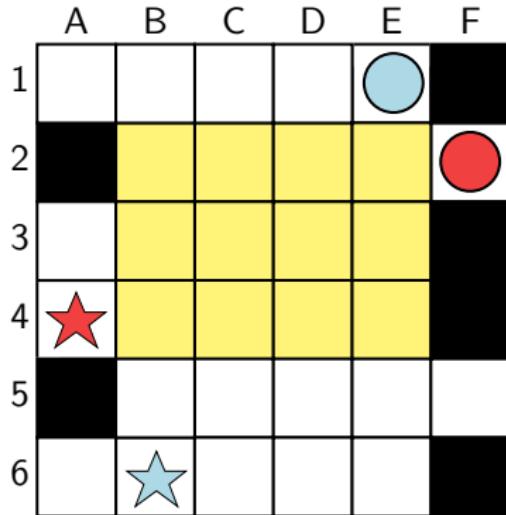
<https://bojie-shen.com/>

# Rectangle Symmetry



Each agent has many equivalent paths through the yellow area. Solving this problem requires proving that one agent must wait.

# Rectangle Symmetry



R	1	2	3	4	5	6	7	8	9
1	1	1	2	3	4	5	6	7	8
2		3	7	14	26	46	79	133	221
3			22	53	116	239	472	904	1,692
4				142	392	1,016	2,651	6,828	17,747
5					1,015	2,971	8,525	23,733	65,236
6						7,447	24,275	78,002	254,173
7							62,429	222,524	795,197
8								573,004	>1,518,151

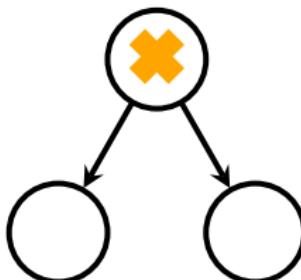
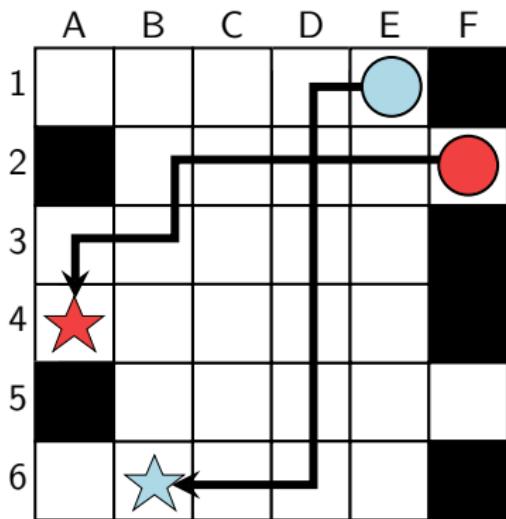
## Table:

Nodes expanded by CBS to solve a 2-agent problem with a  $n \times m$  rectangle.

Each agent has many equivalent paths through the yellow area. Solving this problem requires proving that one agent must wait.

# Barrier Constraints

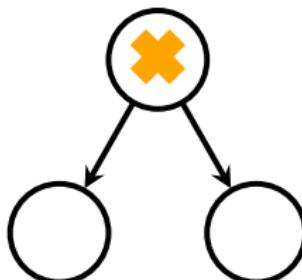
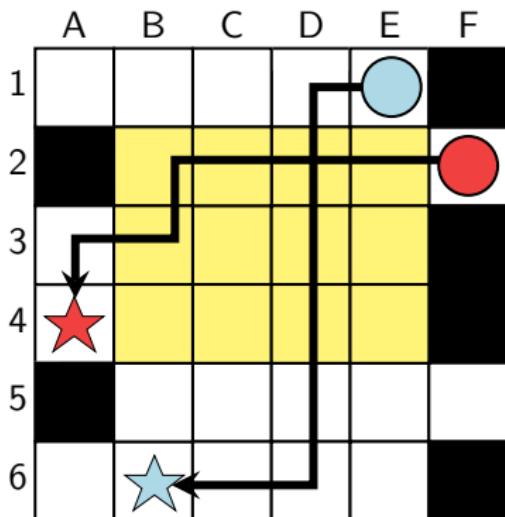
Jiaoyang Li, Daniel Harabor, Peter J. Stuckey, Hang Ma, Sven Koenig. **Symmetry Breaking Constraints for Grid-based Multi-Agent Path Finding.** AAAI, 2019.



Barriers constraints are a specialised reasoning technique that *break rectangle symmetries* and help CBS prove optimality sooner.

# Barrier Constraints

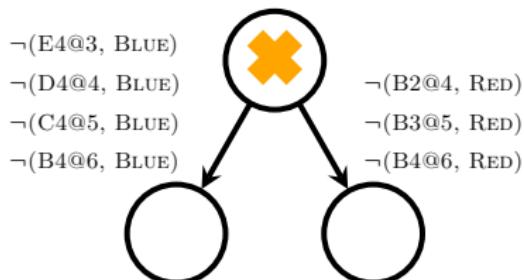
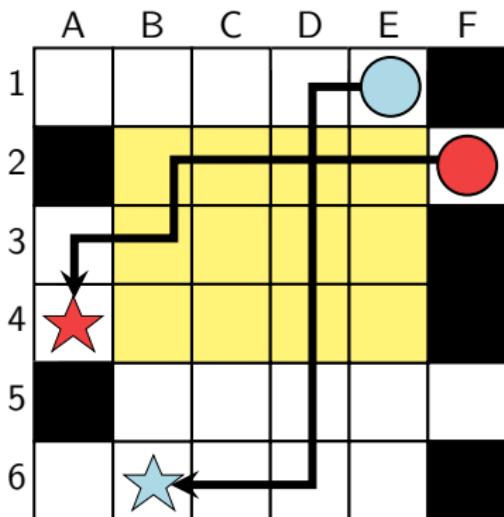
Jiaoyang Li, Daniel Harabor, Peter J. Stuckey, Hang Ma, Sven Koenig. **Symmetry Breaking Constraints for Grid-based Multi-Agent Path Finding**. AAAI, 2019.



We modify CBS to look for conflicts which have rectangle symmetries.

# Barrier Constraints

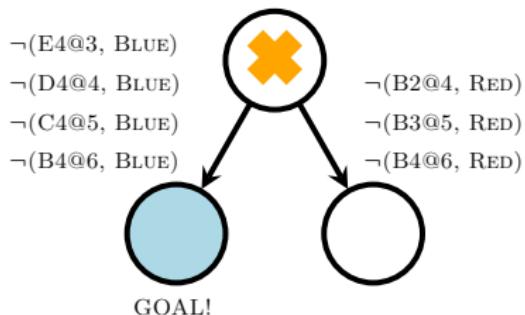
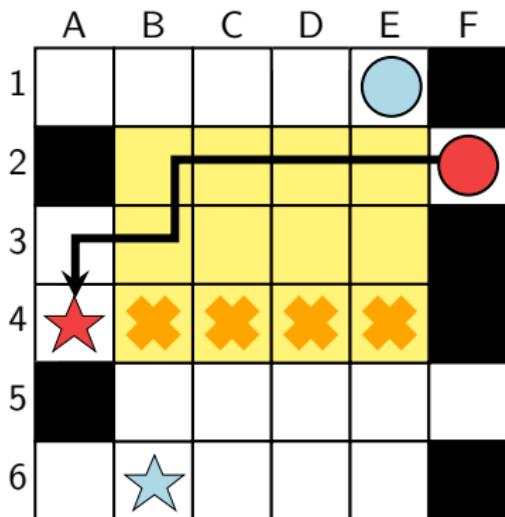
Jiaoyang Li, Daniel Harabor, Peter J. Stuckey, Hang Ma, Sven Koenig. **Symmetry Breaking Constraints for Grid-based Multi-Agent Path Finding.** AAAI, 2019.



When we detect a conflict with rectangle symmetries we post a set of constraints that force agents to *wait* or find a *detour*.

# Barrier Constraints

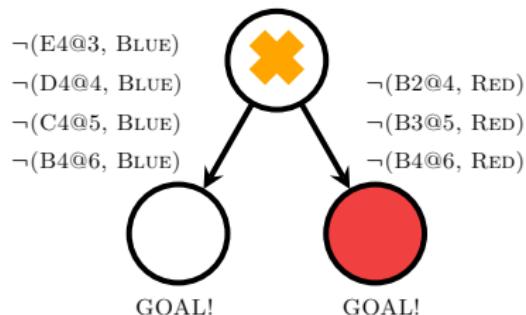
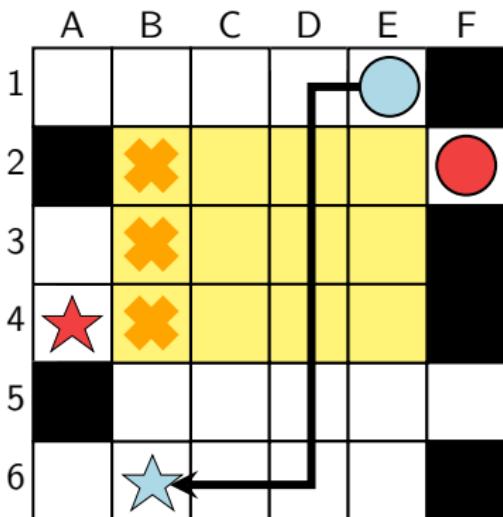
Jiaoyang Li, Daniel Harabor, Peter J. Stuckey, Hang Ma, Sven Koenig. **Symmetry Breaking Constraints for Grid-based Multi-Agent Path Finding.** AAAI, 2019.



We replan Blue. Here, the barrier forces the agent to wait. Blue can wait anywhere but we prefer a path that avoids Red.

# Barrier Constraints

Jiaoyang Li, Daniel Harabor, Peter J. Stuckey, Hang Ma, Sven Koenig. **Symmetry Breaking Constraints for Grid-based Multi-Agent Path Finding.** AAAI, 2019.



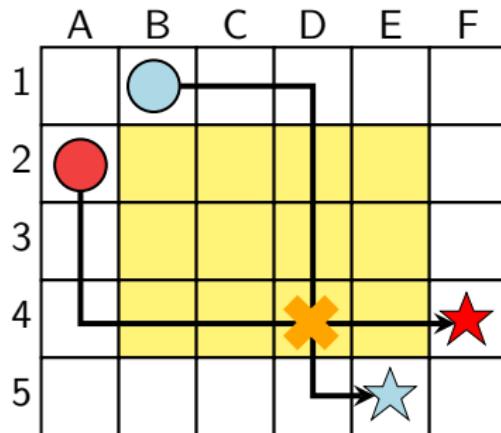
We replan Red. Again, the barrier forces the agent to wait. We can choose any optimal path but prefer one that avoids Blue.

## Cardinal Rectangle Reasoning (CR)

Each rectangle so far satisfies the following conditions:

- (a) Both agents have a Manhattan optimal path, from start to target.
- (b) Both agents reach every tile inside the rectangle area, at the same optimal time.
- (c) All optimal paths pass through the rectangle area and are in collision.

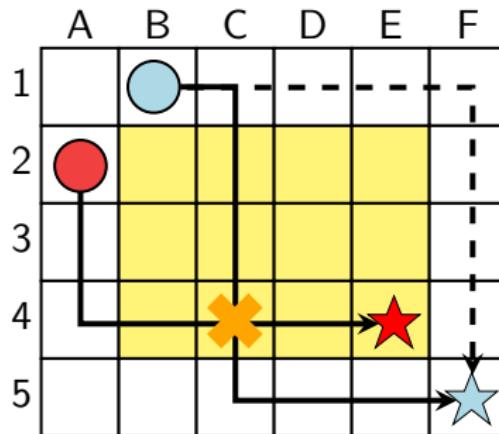
# Cardinal Rectangle Reasoning (CR)



Resolving such *cardinal rectangles* increases solution cost by 1.

# Rectangle Reasoning for Paths (R)

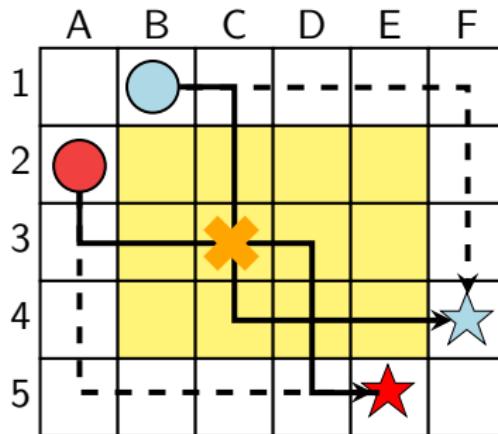
Relaxing the three conditions allows us to detect other types of rectangle symmetries.



Resolving this *semi-cardinal* rectangle raises the cost for only one agent.

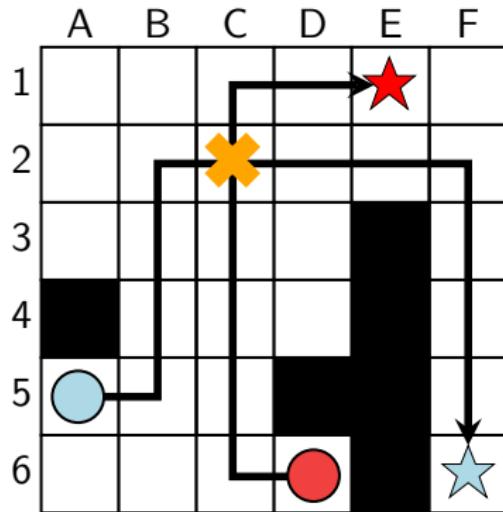
## Rectangle Reasoning for Paths (R)

Relaxing the three conditions allows us to detect other types of rectangle symmetries.



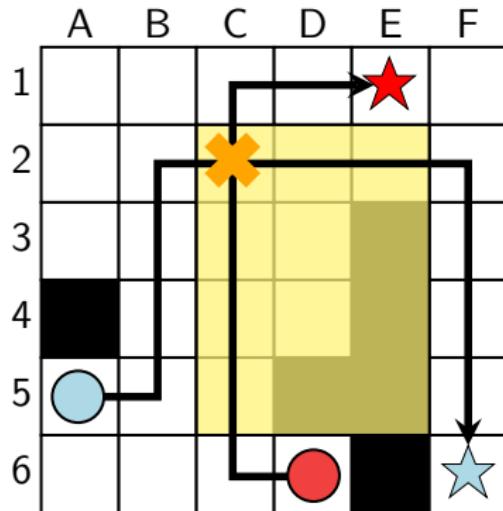
Resolving this *non-cardinal* rectangle eliminates some optimal paths but does not raise any costs.

# Rectangle Reasoning for Path Segments (RM)



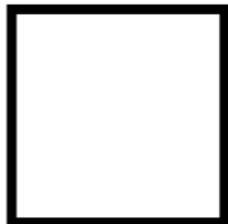
Neither of these agents has a Manhattan optimal path to their target.

# Rectangle Reasoning for Path Segments (RM)

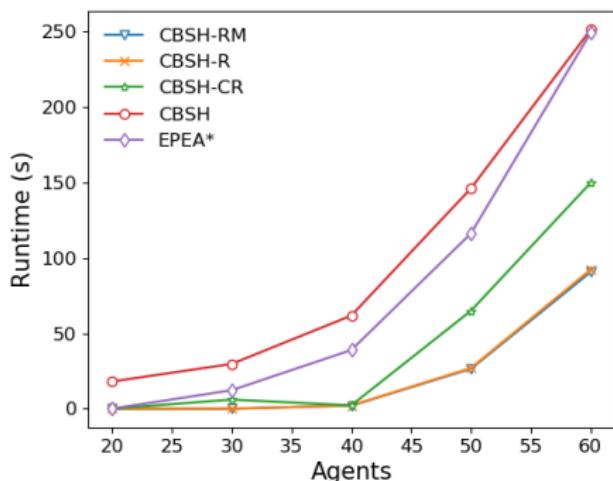
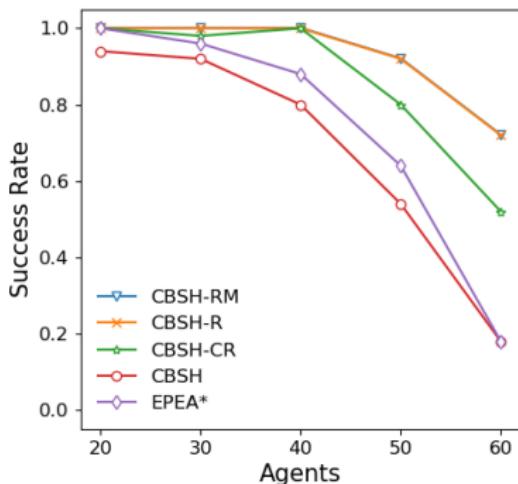


Rectangle symmetries exist, but only for Manhattan optimal subpaths.

# Experiments on Small Grids



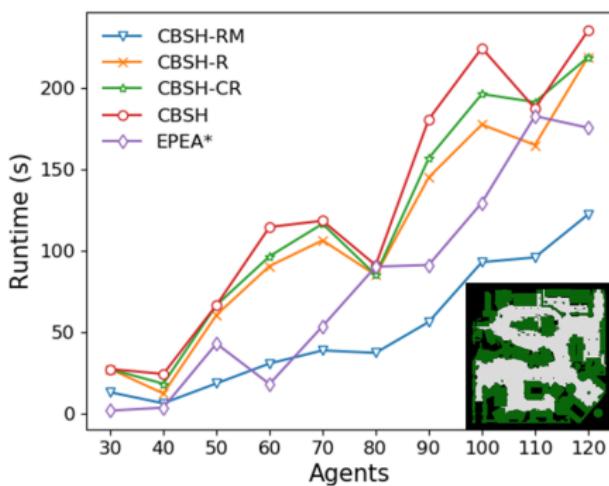
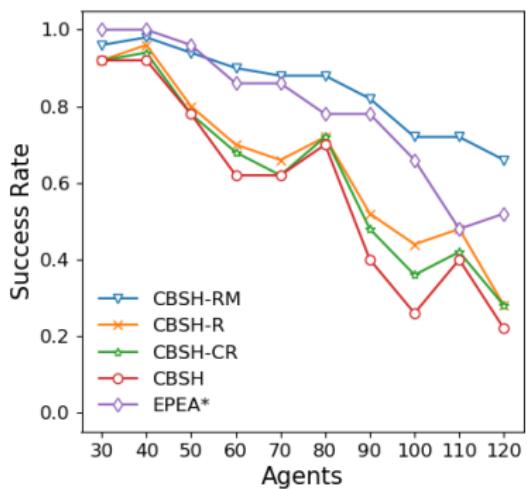
- ▶ Random start-target locations
- ▶ 50 instances per problem size (# agents).
- ▶ Sum-of-Costs objective
- ▶ 20x20 empty grid



# Experiments on Large Grids



- ▶ Random start-target locations
- ▶ 50 instances per problem size (# agents).
- ▶ Sum-of-Costs objective
- ▶ *den520d*:  $194 \times 194$  game grid.



# Corridor Symmetry

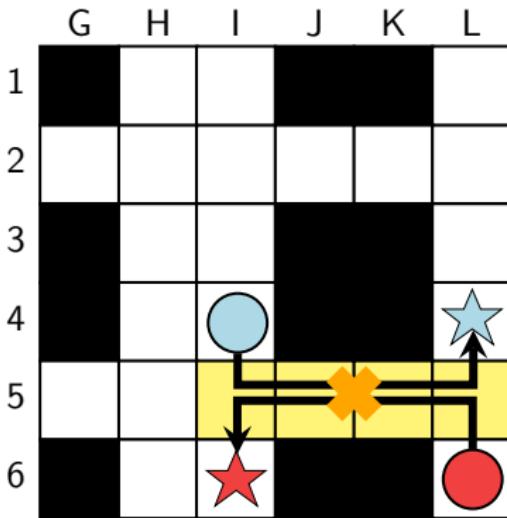
Li, J., Gange, G., Harabor, D., Stuckey, P.J., Ma, H. and Koenig, S., 2020. **New techniques for pairwise symmetry breaking in multi-agent path finding.** In Proceedings of ICAPS.



Corridor symmetries arise when replanning agents after a head-to-head collision in a congested area.

# Corridor Symmetry

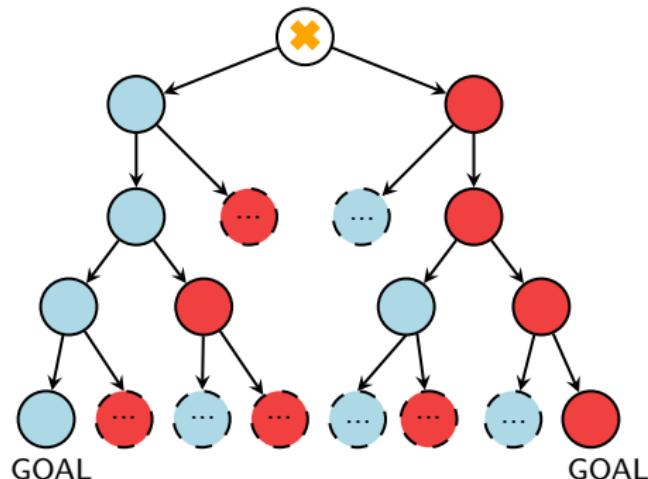
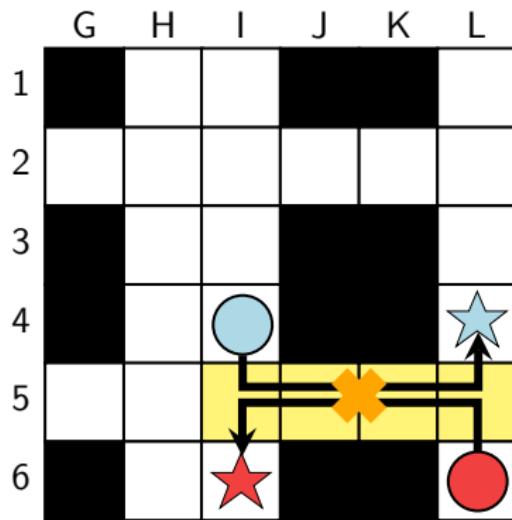
Li, J., Gange, G., Harabor, D., Stuckey, P.J., Ma, H. and Koenig, S., 2020. **New techniques for pairwise symmetry breaking in multi-agent path finding**. In Proceedings of ICAPS.



The optimal strategy is for one of the agents to postpone its corridor traversal (or to find an optimal-cost detour).

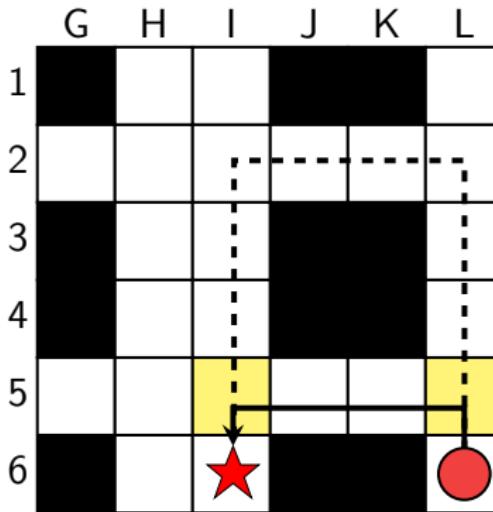
# Corridor Symmetry

Li, J., Gange, G., Harabor, D., Stuckey, P.J., Ma, H. and Koenig, S., 2020. **New techniques for pairwise symmetry breaking in multi-agent path finding.** In Proceedings of ICAPS.



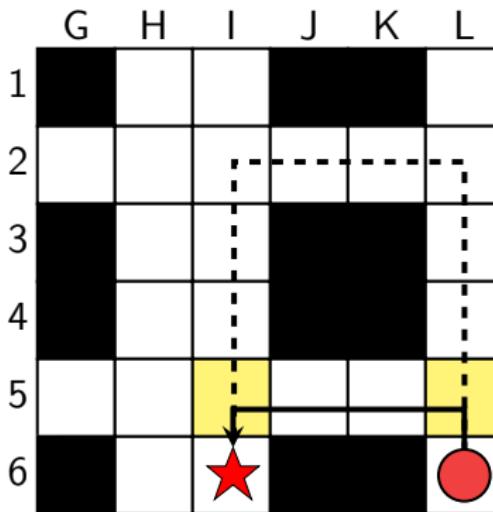
When corridor symmetries arise, they can be very difficult to resolve without specialised reasoning.

# Resolving Corridor Symmetries



Here, red could wait for blue or bypass the corridor.

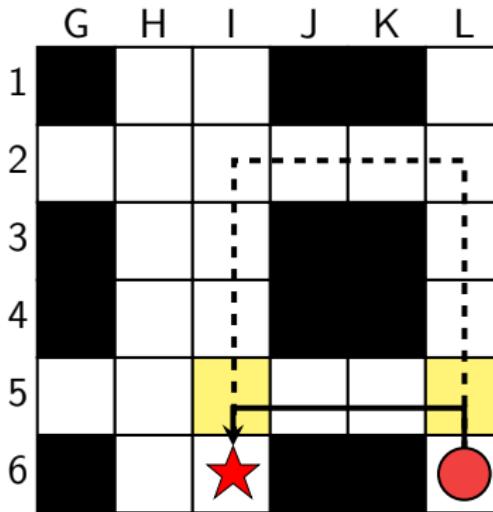
# Resolving Corridor Symmetries



## Wait Option

- ▶ The corridor has length  $k = 3$
- ▶ Blue exits corridor at location L5 at time  $t_2 = 4$ .
- ▶ Red exits corridor at location I5 at time  $t_1 = t_2 + 1 + k = 8$

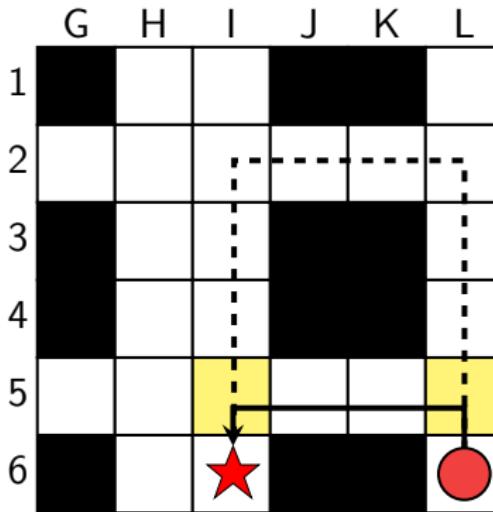
# Resolving Corridor Symmetries



## Bypass Option

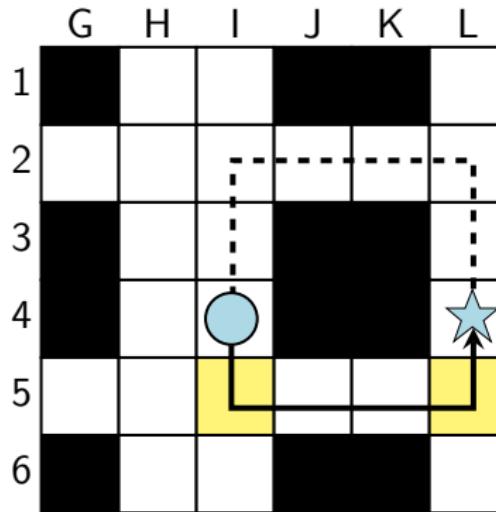
- ▶ Red reaches location I5 at time  $t'_1 = 10$

# Resolving Corridor Symmetries



Waiting looks better, so we add a new constraint:  
 $\langle \text{red}, \text{I}5, [0, \min\{8, 10\}] \rangle$

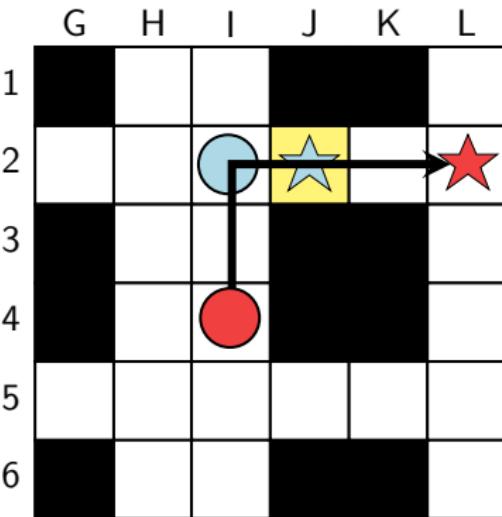
# Resolving Corridor Symmetries



When replanning blue, the bypass option looks better:  
 $\langle \text{blue}, \text{L5}, [0, \min\{9, 7\}] \rangle$

# Target Symmetry

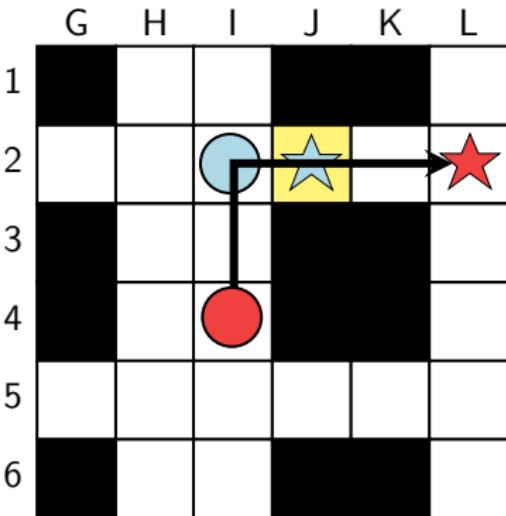
Li, J., Gange, G., Harabor, D., Stuckey, P.J., Ma, H. and Koenig, S., 2020. **New techniques for pairwise symmetry breaking in multi-agent path finding.** In Proceedings of ICAPS.



Target symmetries arise when one agent arrives at its target but later needs to make room for another agent.

# Target Symmetry

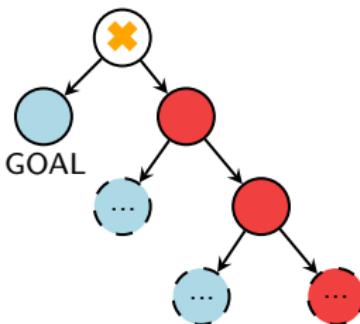
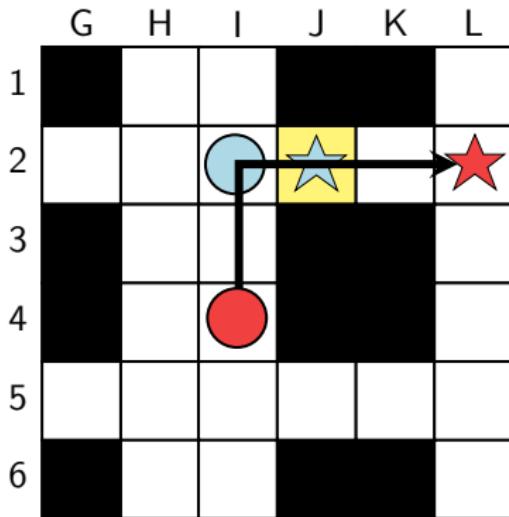
Li, J., Gange, G., Harabor, D., Stuckey, P.J., Ma, H. and Koenig, S., 2020. **New techniques for pairwise symmetry breaking in multi-agent path finding.** In Proceedings of ICAPS.



The optimal strategy is for the first agent to avoid waiting at its target (or for the second to find an optimal-cost detour).

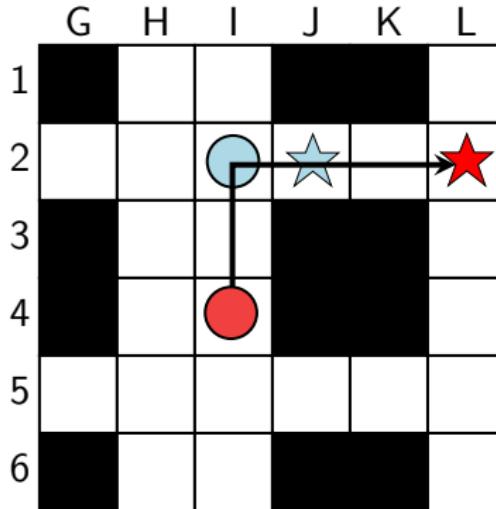
# Target Symmetry

Li, J., Gange, G., Harabor, D., Stuckey, P.J., Ma, H. and Koenig, S., 2020. **New techniques for pairwise symmetry breaking in multi-agent path finding.** In Proceedings of ICAPS.



When target symmetries arise, they also can be difficult to resolve without specialised reasoning.

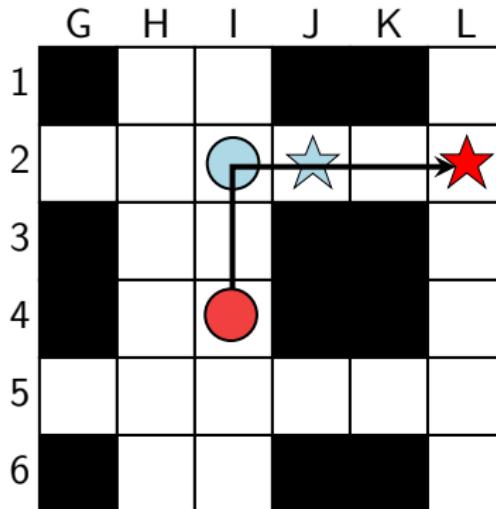
# Resolving Target Symmetries



To resolve the conflict we *split only on blue*.

There are two choices: blue gets priority or blue waits.

# Resolving Target Symmetries

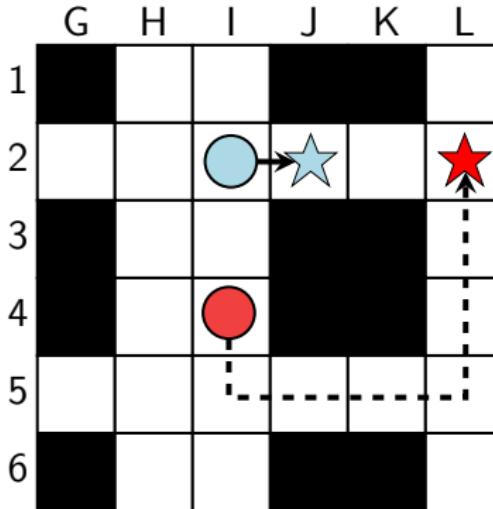


## Blue Waits

We add a constraint on the length of the path for blue:  $l_{\text{blue}} > 3$

Blue finds a new conflict-free path:  $\langle I2, I1, I1, I2, J2 \rangle$

# Resolving Target Symmetries

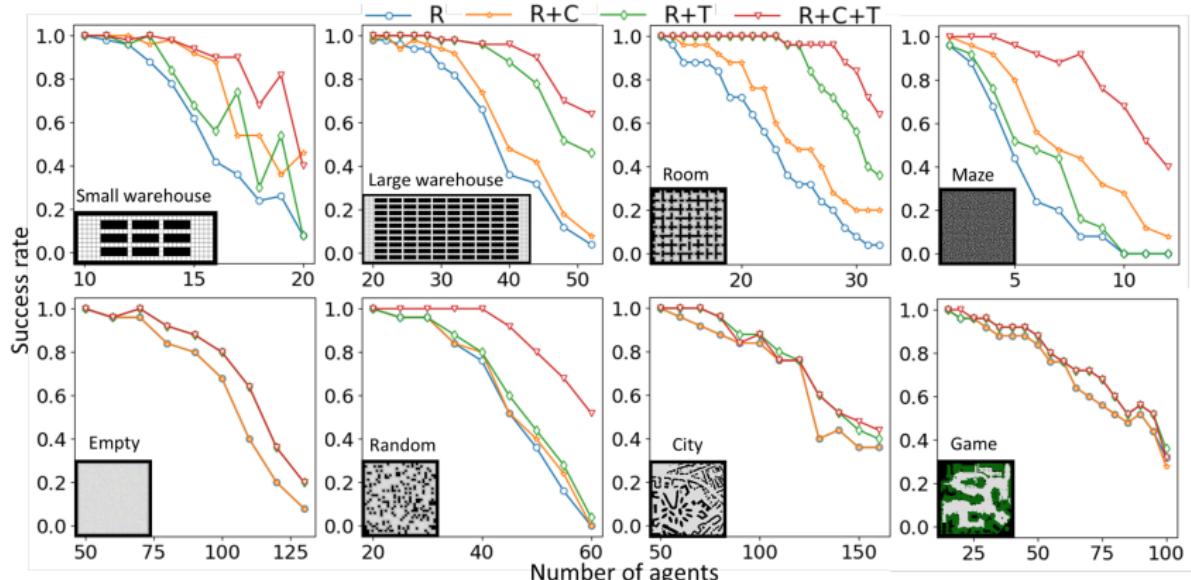


## Blue Gets Priority

We add a constraint on the length of the path for blue:  $l_{\text{blue}} \leq 3$

This forces red to find a bypass.

# Experiments with Corridor and Target Reasoning



## Symmetry Breaking

- ▶ Automatic symmetry detection
  - ▶ Zhang, H., Li, J., Surynek, P., Kumar, T.K.S. and Koenig, S., 2022. Multi-agent path finding with mutex propagation. Artificial Intelligence.

## High-Level Heuristics

- ▶ Pairwise reasoning
  - ▶ Felner, A., Li, J., Boyarski, E., Ma, H., Cohen, L., Kumar, T.K.S. and Koenig, S., 2018. Adding heuristics to conflict-based search for multi-agent path finding. In Proceedings of ICAPS.
  - ▶ Li, J., Felner, A., Boyarski, E., Ma, H. and Koenig, S., 2019. Improved heuristics for multi-agent path finding with conflict-based search. In Proceedings of IJCAI.
- ▶ Beyond pairwise reasoning
  - ▶ Shen, B., Chen, Z., Li, J., Cheema, M.A., Harabor, D. and Stuckey, P.J., 2023. Beyond pairwise reasoning in multi-agent path finding. In Proceedings of ICAPS.

# Awesome MAPF Tools

**Tracker**

Browse

- Home
- Benchmarks**
- Submissions

Make a submission

- Call for submissions
- New submission request
- Submissions and API keys

Docs

- Docs
- About
- Github

Settings

- Light mode
- Hide tips

More >

**Benchmarks**

Map count: 33  
Instance count: 1460645  
Scenario count: 1650  
Instances solved: 100.00%  
Instances closed: 15.44%

References:  
Roni Stern, Nathan Sturtevant, Ariel Felner, Sven Koenig, Hang Ma, Thayne Walker, Jaoyang Li et al. "Multi-agent pathfinding: Definitions, variants, and benchmarks." In Proceedings of the International Symposium on Combinatorial Search (SoCS), 2019.

Browse maps Trends Compare

Map	Size	Domain	Instances Solved	Instances Closed
Berlin 1256	256x256	City	100%	26%
Paris 1256	256x256	City	100%	23%
Boston 0 256	256x256	City	100%	17%
Empty 32 32	32x32	Empty	100%	36%
Empty 8 8	8x8	Empty	100%	98%
Empty 48 48	48x48	Empty	100%	28%
Empty 16 16	16x16	Empty	100%	56%
Lak 303 D	194x194	Game	100%	1-33 of 33

Please visit the website: <https://tracker.pathfinding.ai/>

End

That's it for this week